

# Advanced BPF Kernel Features for the Container Age

Daniel Borkmann, Cilium.io & co-maintainer BPF

FOSDEM 2021





# **BPF, a general purpose engine**

BPF, a general purpose execution engine



Recap: minimal instruction set architecture  
with two major design goals:

BPF, a general purpose execution engine



- 1) Low overhead when mapping to native code, in particular: x86-64, arm64

BPF, a general purpose execution engine



2) Must be verifiable for **safety** by the kernel at program load time.

BPF, a general purpose execution engine



Safety to the extent that developers choose implementing programs in BPF rather than traditional kernel modules.

BPF, a general purpose execution engine



Kernel modules may cause kernel to panic. BPF programs are safety-checked and will **not** crash the kernel.

BPF, a general purpose execution engine



Kernel provides framework of building blocks and attachment points.



BPF, a general purpose execution engine



Is BPF a generic virtual machine? No.

BPF, a general purpose execution engine



Is BPF a fully generic instruction set? No.

BPF, a general purpose execution engine



It's an instruction set with C calling convention in mind.

BPF, a general purpose execution engine



Why? Kernel is written in C and BPF needs  
to **efficiently** interact with the kernel.  
Approx 150 BPF kernel helpers, 30 maps.

BPF, a general purpose execution engine



How far off is “BPF” C from generic C?

BPF, a general purpose execution engine



BPF has seen major advances such as  
BPF-2-BPF function calls, bounded loops,  
global variables, static linking, BTF,  
up to 1 Mio instructions / program, ...

BPF, a general purpose execution engine



This already allows for solving a lot of  
interesting production issues.

Enter: **Cilium**



# **Cilium: cloud-native networking, security & observability built upon BPF**





# Cilium: use cases overview

## Networking

- Highly efficient and flexible networking
- Routing, overlay, cloud-provider native
- IPv4, IPv6, NAT46
- Multi cluster routing

## Load Balancing:

- Highly scalable L3-L4 (XDP) load balancing
- Kubernetes services (replaces kube-proxy)
- Multi-cluster
- Service affinity (prefer zones)

## Network Security

- Identity-based network security
- API-aware security (HTTP, gRPC, ...), DNS-aware
- Transparent encryption

## Observability

- Metrics (Network, DNS, Security, Latencies, HTTP, ...)
- Flow logs (with datapath aggregation)

## ServiceMesh:

- Minimized overhead when injecting servicemesh sidecar proxies
- Istio integration



Recent Cilium 1.9 and BPF kernel extensions

# Deep dive 1: Kubernetes service load balancing with XDP, BPF & Maglev



## Recent Cilium 1.9 and BPF kernel extensions

kube-proxy example: co-location of service load balancer with regular user workloads on every node.




## Recent Cilium 1.9 and BPF kernel extensions

Migrating (N-S) load balancing from legacy subsystems to BPF at XDP layer reduces CPU cost significantly, and achieves DPDK speeds.




# Recent Cilium 1.9 and BPF kernel extensions





# Recent Cilium 1.9 and BPF kernel extensions






## Recent Cilium 1.9 and BPF kernel extensions

... all with BPF on upstream kernel drivers.  
No busy-polling CPUs. No user-kernel  
boundary crossing.




# K8s LoadBalancer service on-prem example








# K8s LoadBalancer service on-prem example





# K8s LoadBalancer service on-prem example





## K8s LoadBalancer service on-prem example

LoadBalancer implementation done by  
Cloud providers or MetalLB for on-prem.  
MetalLB can announce via ARP/NDP or BGP.



## K8s LoadBalancer service on-prem example

MetalLB does IP address allocation and external announcement, but does not sit in critical fast path (hence works with XDP).





## K8s LoadBalancer service on-prem example

Default backend selection random with  
node-local BPF CT for remaining sticky to  
backend. However ...




# Recent Cilium 1.9 and BPF kernel extensions

Random




backend






# Recent Cilium 1.9 and BPF kernel extensions

Random






# Recent Cilium 1.9 and BPF kernel extensions








# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions





# Recent Cilium 1.9 and BPF kernel extensions







## Recent Cilium 1.9 and BPF kernel extensions

Needs consistent hashing for resiliency against failures. Cilium 1.9 supports Maglev-based selection from BPF/XDP side.



# Recent Cilium 1.9 and BPF kernel extensions

## Maglev




**backend**





# Recent Cilium 1.9 and BPF kernel extensions


## Maglev








# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions





# Recent Cilium 1.9 and BPF kernel extensions





## Recent Cilium 1.9 and BPF kernel extensions

BPF/XDP datapath performs dynamic sized map-in-map lookup for optimized memory utilization. Backend selected via jhash of tuple over per-service Maglev table.

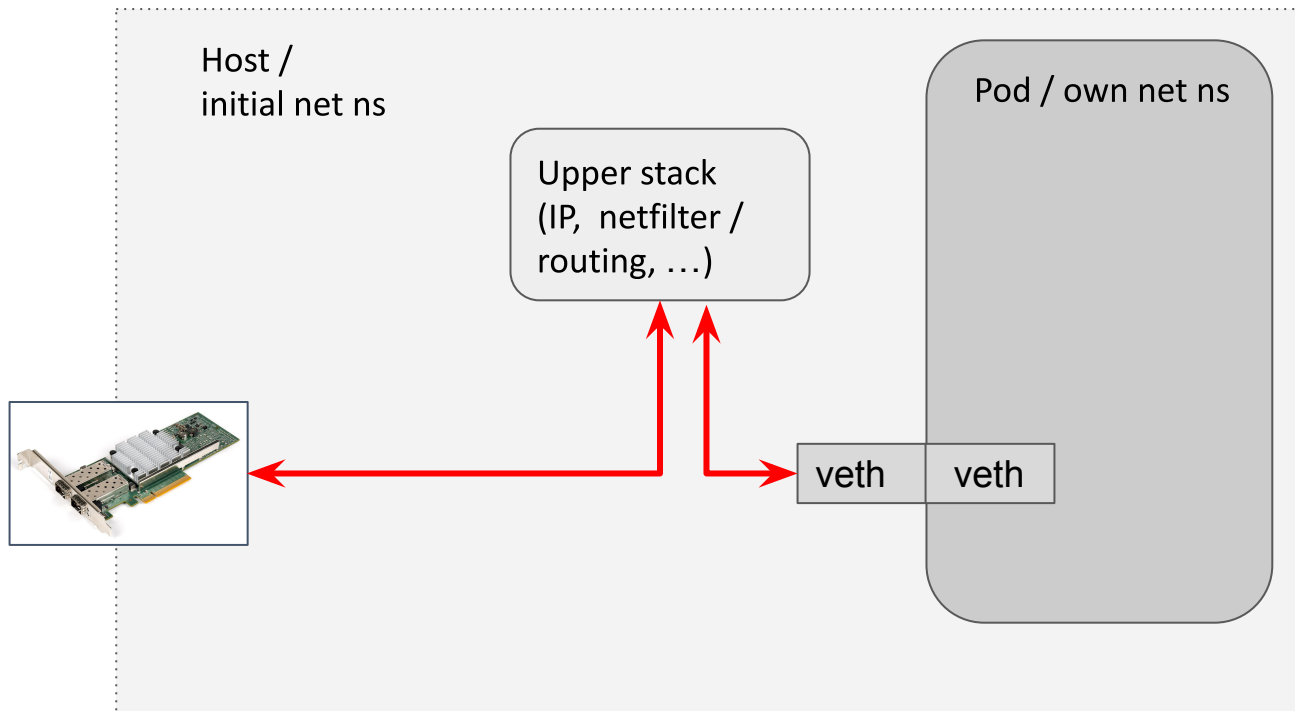


Recent Cilium 1.9 and BPF kernel extensions

## Deep dive 2: low-latency Pod fast-path with BPF




# Recent Cilium 1.9 and BPF kernel extensions







# Recent Cilium 1.9 and BPF kernel extensions






# Recent Cilium 1.9 and BPF kernel extensions

High-level internals:

```
dev = ops->ndo_get_peer_dev(dev)
skb_scrub_packet()
skb->dev = dev
sch_handle_ingress():
- goto another_round
- no CPU backlog queue
```

Upper stack  
(IP, netfilter /  
routing, ...)

bpf\_redirect\_peer()



veth

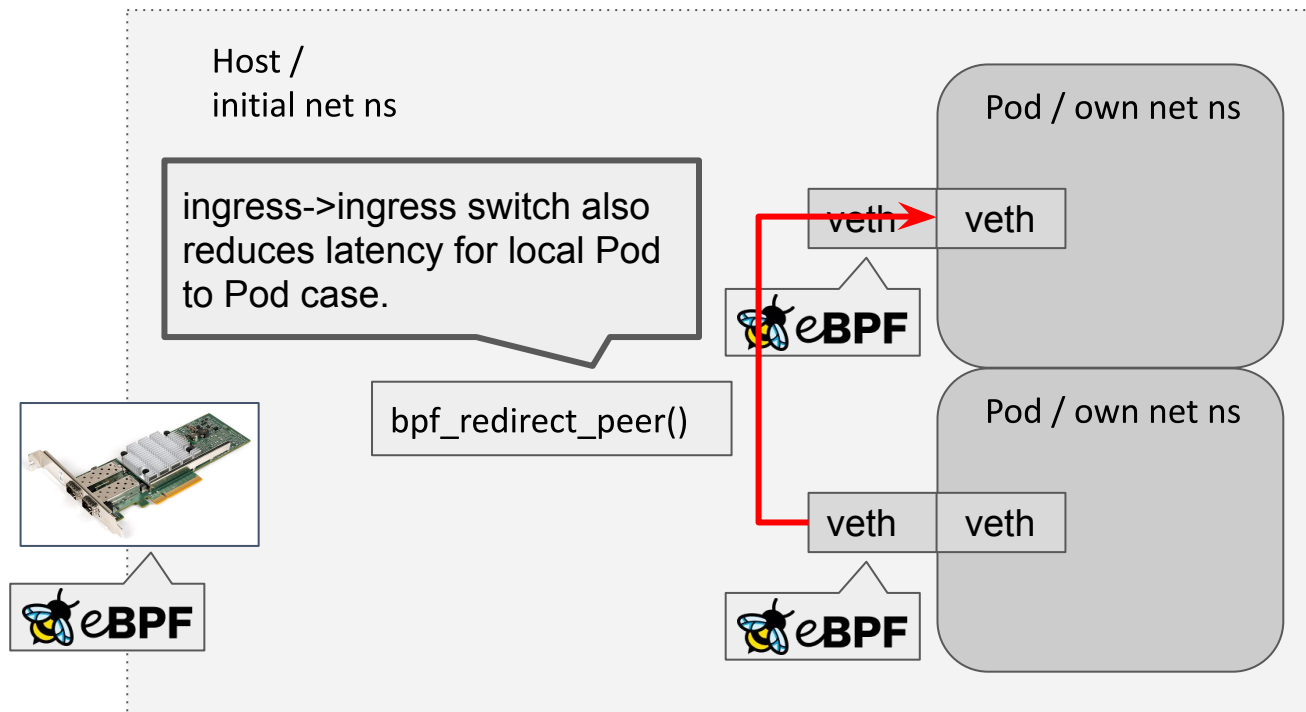
veth



Pod / own net ns




# Recent Cilium 1.9 and BPF kernel extensions





# Recent Cilium 1.9 and BPF kernel extensions





## Recent Cilium 1.9 and BPF kernel extensions

Low-latency net ns switch into Pod from BPF.  
Automatic L2 resolution for traffic from Pod.  
Both directions now avoid host stack.


<https://git.kernel.org/torvalds/c/9aa1206e8f48>

<https://git.kernel.org/torvalds/c/b4ab31414970>



# Recent Cilium 1.9 and BPF kernel extensions

TCP\_SENDFILE performance, single stream, v5.10 (higher is better)




Xeon E3-1240, 3.4 GHz each (4 cores, HT off), back to back via nfp, IRQs pinned to CPUs, tuned with profile network-throughput  
From Pod ns: netperf -H <remote-pod-ip> -t TCP\_SENDFILE -T0,0 -P0 -s2 -l 60 -D 2 -f g



# Recent Cilium 1.9 and BPF kernel extensions

TCP\_RR performance, single session, v5.10 (higher is better)



Xeon E3-1240, 3.4 GHz each (4 cores, HT off), back to back via nfp, IRQs pinned to CPUs, tuned with profile network-latency  
From Pod ns: percpu\_netperf <remote-pod-ip> ([https://github.com/borkmann/netperf\\_scripts/blob/master/percpu\\_netperf](https://github.com/borkmann/netperf_scripts/blob/master/percpu_netperf))



Recent Cilium 1.9 and BPF kernel extensions

## Deep dive 3: EDT rate-limiting for Pods via BPF





## Recent Cilium 1.9 and BPF kernel extensions

Old way: CNI chaining with bandwidth plugin that sets up TBF qdisc. Even sets up ifb for ingress shaping. Not scalable.



Recent Cilium 1.9 and BPF kernel extensions

Cilium native: lock-less Pod rate-limiting on multi-queue via BPF and Earliest Departure Time (EDT).




## Recent Cilium 1.9 and BPF kernel extensions

BPF classifies network traffic to Pod and then sets packet departure time (`skb->tstamp`) based on user defined bandwidth rate.



## Recent Cilium 1.9 and BPF kernel extensions


FQ qdisc schedules packet under this  
time constraint.





# Recent Cilium 1.9 and BPF kernel extensions

**Transmission latency (lower is better)**



**99p → 10x latency reduction**

**95p → 20x latency reduction**

**50p**

## Recap



Unlike other kernel subsystems BPF is in a unique position to innovatively tackle and optimize such complex production issues.

## Recap



Cilium with BPF at its core delivers this technological revolution to mainstream Kubernetes.





**Urs Hölzle** @uhoelzle · Aug 29

ICYMI: GKE is, once again, the first managed kubernetes service to provide a key new feature. This time it's better networking with **eBPF** support.



**Tim Hockin**  
@thockin

Cilium and the team around it have impressed me from the beginning. My mind is spinning with the possibilities of eBPF.



**Kelsey Hightower** ✓  
@kelseyhightower

The latest @ciliumproject release is taking this Kubernetes networking thing to another level: multi-cluster service routing, IPVLAN support, transparent encryption, and DNS authorization.

[cilium.io/blog/2019/02/1...](https://cilium.io/blog/2019/02/1...)

# Thanks! Questions?



[github.com/cilium/cilium](https://github.com/cilium/cilium)

[cilium.io](https://cilium.io)

[ebpf.io](https://ebpf.io)

